

lab-4-perceptron-adaline-torchDiff

March 4, 2023

```
[182]: import numpy as np
import pandas as pd
import torch
import matplotlib.pyplot as plt
```

0.1 Implementing numpy version of perceptron

```
[263]: data = np.genfromtxt('data/perceptron_toydata.txt', delimiter='\t')
X, y = data[:, :2], data[:, 2]
y = y.astype(int)

print('Class label counts:', np.bincount(y))
print('X.shape:', X.shape)
print('y.shape:', y.shape)

# Shuffling & train/test split
shuffle_idx = np.arange(y.shape[0])
shuffle_rng = np.random.RandomState(123)
shuffle_rng.shuffle(shuffle_idx)
X, y = X[shuffle_idx], y[shuffle_idx]

X_train, X_test = X[shuffle_idx[:70]], X[shuffle_idx[70:]]
y_train, y_test = y[shuffle_idx[:70]], y[shuffle_idx[70:]]

# Normalize (mean zero, unit variance)
mu, sigma = X_train.mean(axis=0), X_train.std(axis=0)
X_train = (X_train - mu) / sigma
X_test = (X_test - mu) / sigma
```

```
Class label counts: [50 50]
X.shape: (100, 2)
y.shape: (100,)
```


```
[19]: # After stadardisation
print("Mean = {}".format(np.mean(X_train, axis=0)))
print("STD = {}".format(np.std(X_train, axis=0)))
```

```
Mean = [2.06184276e-17 7.93016446e-18]
STD = [1. 1.]
```

```
[24]: X_train.shape
```


```
[24]: (70, 2)
```

```
[28]: plt.scatter(X_train[y_train==0, 0], X_train[y_train==0, 1], label='class 0',  
                 ↪marker='o')  
plt.scatter(X_train[y_train==1, 0], X_train[y_train==1, 1], label='class 1',  
                 ↪marker='s')  
plt.title('Training set')  
plt.xlabel('feature 1')  
plt.ylabel('feature 2')  
plt.xlim([-3, 3])  
plt.ylim([-3, 3])  
plt.legend()  
plt.show()
```



```
[32]: plt.scatter(X_test[y_test==0, 0], X_test[y_test==0, 1], label='class 0',  
                 ↪marker='o')  
plt.scatter(X_test[y_test==1, 0], X_test[y_test==1, 1], label='class 1',  
                 ↪marker='s')
```

```
plt.title('Test set')
plt.xlabel('feature 1')
plt.ylabel('feature 2')
plt.xlim([-3, 3])
plt.ylim([-3, 3])
plt.legend()
plt.show()
```



```
[176]: class Perceptron():
    def __init__(self, num_features):
        self.num_features = num_features
        self.weights = np.zeros((num_features,1), dtype=float)
        self.bias = np.zeros(1, dtype=float)

    def forward(self, x):
        linear = np.dot(x, self.weights) + self.bias
        linear = linear.flatten()
        predictions = np.where(linear>0,1,0)
        return predictions

    def backward(self, x,y):
        predictions = self.forward(x)
        errors = y - predictions
```

```

        return errors

    def newTrain(self, x,y,epochs):
        for e in range(epochs):
            error = self.backward(x,y).reshape(x.shape[0],1)
            error = error*x
            error = np.sum(error, axis=0).reshape(-1,1)
            self.weights+=error
            self.bias+=error[0,0]

    def train(self, x,y,epochs):
        for e in range(epochs):
            for i in range(y.shape[0]):
                error = self.backward(x[i].reshape(1, self.num_features),y[i]).
↪reshape(-1)
                self.weights += (error*x[i]).reshape(self.num_features,1)
                self.bias += error

    def evaluate(self,x,y):
        predictions = self.forward(x).reshape(-1)
        accuracy = np.sum(predictions == y) / y.shape[0]
        return accuracy

```

```

[177]: ppn = Perceptron(num_features=2)

ppn.newTrain(X_train, y_train, epochs=1)

print('Model parameters:\n\n')
print('  Weights: %s\n' % ppn.weights)
print('  Bias: %s\n' % ppn.bias)

```

(2, 1)

Model parameters:

Weights: [[29.83124073]
[28.73135694]]

Bias: [29.83124073]

```

[179]: ppn.evaluate(X_test, y_test)*100

```

[179]: 93.33333333333333

0.2 Fully Connected Linear layer

```
[61]: X = torch.arange(50, dtype=torch.float32).view(10,5)

[65]: fc_layer = torch.nn.Linear(in_features=5, out_features = 3)

[66]: fc_layer.weight.shape

[66]: torch.Size([3, 5])

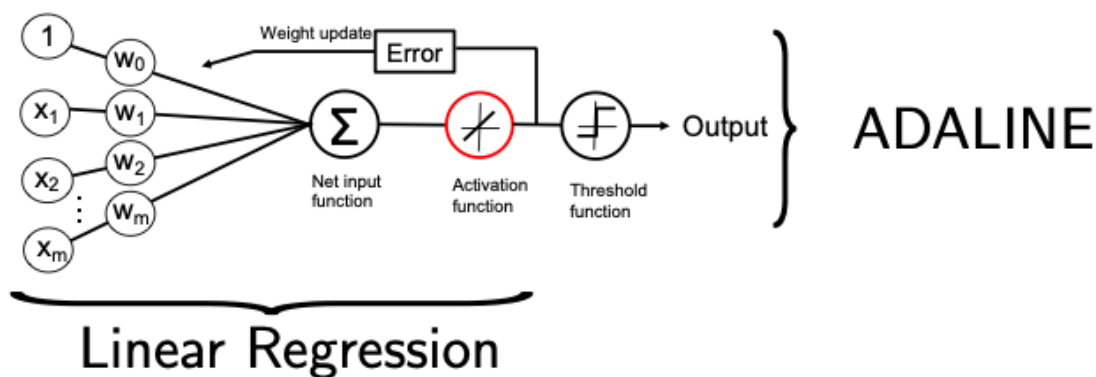
[67]: fc_layer.bias.shape

[67]: torch.Size([3])

[68]: fc_layer(X).shape

[68]: torch.Size([10, 3])
```

0.3 Adaline



```
[269]: # Load the data
df = pd.read_csv('data/linreg-data.csv', index_col=0)
df.tail()

[269]:
```

| | x1 | x2 | y |
|-----|-----------|-----------|------------|
| 995 | -0.942094 | -0.835856 | -22.324428 |
| 996 | 1.222445 | -0.403177 | -52.121493 |
| 997 | -0.112466 | -1.688230 | -57.043196 |
| 998 | -0.403459 | -0.412272 | -27.701833 |
| 999 | 0.021351 | -0.499017 | -9.804714 |

```


[270]: X = torch.tensor(df[['x1', 'x2']].values, dtype=torch.float)
y = torch.tensor(df['y'].values, dtype=torch.float)

shuffle_idx = torch.randperm(y.size(0), dtype=torch.long)
```

```

X, y = X[shuffle_idx], y[shuffle_idx]

percent70 = int(shuffle_idx.size(0)*0.7)

X_train, X_test = X[shuffle_idx[:percent70]], X[shuffle_idx[percent70:]]
y_train, y_test = y[shuffle_idx[:percent70]], y[shuffle_idx[percent70:]]

# Normalize (mean zero, unit variance)

mu, sigma = X_train.mean(dim=0), X_train.std(dim=0)
X_train = (X_train - mu) / sigma
X_test = (X_test - mu) / sigma

```

```

[193]: class LinearRegression1():
    def __init__(self, num_features):
        self.num_features = num_features
        self.weights = torch.zeros(num_features, 1, dtype=torch.float)
        self.bias = torch.zeros(1, dtype=torch.float)

    def forward(self, x):
        netinputs = torch.add(torch.mm(x, self.weights), self.bias)
        activations = netinputs
        return activations.view(-1)

    def backward(self, x, yhat, y):

        grad_loss_yhat = 2*(yhat - y) #derivative of (yhat-y)**2 wrt yhat.
        ↪ Verified

        grad_yhat_weights = x
        grad_yhat_bias = 1.

        # Chain rule: inner times outer
        grad_loss_weights = torch.mm(grad_yhat_weights.t(), grad_loss_yhat.
        ↪ view(-1, 1)) / y.size(0)

        grad_loss_bias = torch.sum(grad_yhat_bias*grad_loss_yhat) / y.size(0)

        # return negative gradient
        return (-1)*grad_loss_weights, (-1)*grad_loss_bias

```

[194]: *#Define Training and Evaluation Functions*

```

# mean squared error
def loss(yhat, y):
    return torch.mean((yhat - y)**2)

```

```

# batch gradient descent
def train(model, x, y, num_epochs, learning_rate=0.01):
    cost = []
    for e in range(num_epochs):

        ##### Compute outputs #####
        yhat = model.forward(x)

        ##### Compute gradients #####
        negative_grad_w, negative_grad_b = model.backward(x, yhat, y)

        ##### Update weights #####
        model.weights += learning_rate * negative_grad_w
        model.bias += learning_rate * negative_grad_b

        ##### Logging #####
        yhat = model.forward(x) # not that this is a bit wasteful here
        curr_loss = loss(yhat, y)
        print('Epoch: %03d' % (e+1), end="")
        print(' | MSE: %.5f' % curr_loss)
        cost.append(curr_loss)

    return cost

```

```

[195]: model = LinearRegression1(num_features=X_train.size(1))
       cost = train(model, X_train, y_train, num_epochs=100, learning_rate=0.05)

```

```


Epoch: 001 | MSE: 1636.67786
Epoch: 002 | MSE: 1397.68274
Epoch: 003 | MSE: 1204.04407
Epoch: 004 | MSE: 1047.15381
Epoch: 005 | MSE: 920.03729
Epoch: 006 | MSE: 817.04437
Epoch: 007 | MSE: 733.59680
Epoch: 008 | MSE: 665.98505
Epoch: 009 | MSE: 611.20398
Epoch: 010 | MSE: 566.81860
Epoch: 011 | MSE: 530.85614
Epoch: 012 | MSE: 501.71808
Epoch: 013 | MSE: 478.10925
Epoch: 014 | MSE: 458.98059
Epoch: 015 | MSE: 443.48166
Epoch: 016 | MSE: 430.92383
Epoch: 017 | MSE: 420.74896
Epoch: 018 | MSE: 412.50482

```

Epoch: 019 | MSE: 405.82495
Epoch: 020 | MSE: 400.41263
Epoch: 021 | MSE: 396.02731
Epoch: 022 | MSE: 392.47412
Epoch: 023 | MSE: 389.59509
Epoch: 024 | MSE: 387.26233
Epoch: 025 | MSE: 385.37228
Epoch: 026 | MSE: 383.84076
Epoch: 027 | MSE: 382.59982
Epoch: 028 | MSE: 381.59442
Epoch: 029 | MSE: 380.77972
Epoch: 030 | MSE: 380.11963
Epoch: 031 | MSE: 379.58475
Epoch: 032 | MSE: 379.15140
Epoch: 033 | MSE: 378.80023
Epoch: 034 | MSE: 378.51572
Epoch: 035 | MSE: 378.28519
Epoch: 036 | MSE: 378.09836
Epoch: 037 | MSE: 377.94696
Epoch: 038 | MSE: 377.82428
Epoch: 039 | MSE: 377.72501
Epoch: 040 | MSE: 377.64441
Epoch: 041 | MSE: 377.57919
Epoch: 042 | MSE: 377.52631
Epoch: 043 | MSE: 377.48343
Epoch: 044 | MSE: 377.44876
Epoch: 045 | MSE: 377.42062
Epoch: 046 | MSE: 377.39786
Epoch: 047 | MSE: 377.37936
Epoch: 048 | MSE: 377.36438
Epoch: 049 | MSE: 377.35226
Epoch: 050 | MSE: 377.34247
Epoch: 051 | MSE: 377.33441
Epoch: 052 | MSE: 377.32803
Epoch: 053 | MSE: 377.32281
Epoch: 054 | MSE: 377.31857
Epoch: 055 | MSE: 377.31519
Epoch: 056 | MSE: 377.31238
Epoch: 057 | MSE: 377.31009
Epoch: 058 | MSE: 377.30829
Epoch: 059 | MSE: 377.30679
Epoch: 060 | MSE: 377.30557
Epoch: 061 | MSE: 377.30466
Epoch: 062 | MSE: 377.30383
Epoch: 063 | MSE: 377.30322
Epoch: 064 | MSE: 377.30273
Epoch: 065 | MSE: 377.30228
Epoch: 066 | MSE: 377.30191


```
Epoch: 067 | MSE: 377.30167
Epoch: 068 | MSE: 377.30142
Epoch: 069 | MSE: 377.30124
Epoch: 070 | MSE: 377.30112
Epoch: 071 | MSE: 377.30099
Epoch: 072 | MSE: 377.30090
Epoch: 073 | MSE: 377.30081
Epoch: 074 | MSE: 377.30072
Epoch: 075 | MSE: 377.30072
Epoch: 076 | MSE: 377.30066
Epoch: 077 | MSE: 377.30063
Epoch: 078 | MSE: 377.30063
Epoch: 079 | MSE: 377.30057
Epoch: 080 | MSE: 377.30057
Epoch: 081 | MSE: 377.30048
Epoch: 082 | MSE: 377.30054
Epoch: 083 | MSE: 377.30054
Epoch: 084 | MSE: 377.30054
Epoch: 085 | MSE: 377.30054
Epoch: 086 | MSE: 377.30048
Epoch: 087 | MSE: 377.30048
Epoch: 088 | MSE: 377.30048
Epoch: 089 | MSE: 377.30045
Epoch: 090 | MSE: 377.30048
Epoch: 091 | MSE: 377.30045
Epoch: 092 | MSE: 377.30048
Epoch: 093 | MSE: 377.30045
Epoch: 094 | MSE: 377.30048
Epoch: 095 | MSE: 377.30045
Epoch: 096 | MSE: 377.30045
Epoch: 097 | MSE: 377.30048
Epoch: 098 | MSE: 377.30045
Epoch: 099 | MSE: 377.30048
Epoch: 100 | MSE: 377.30048
```

```
[196]: # plot MSE
plt.plot(range(len(cost)), cost)
plt.ylabel('Mean Squared Error')
plt.xlabel('Epoch')
plt.show()
```



```
[197]: train_pred = model.forward(X_train)
       test_pred = model.forward(X_test)

       print('Train MSE: %.5f' % loss(train_pred, y_train))
       print('Test MSE: %.5f' % loss(test_pred, y_test))
```

```
Train MSE: 377.30048
Test MSE: 394.55591
```

Try Modify the `train()` function such that the dataset is shuffled prior to each epoch. Do you see a difference – Yes/No? Try to come up with an explanation for your observation.

0.4 Adaline With SGD

```
[275]: df = pd.read_csv('data/iris.data', index_col=None, header=None)
       df.columns = ['x1', 'x2', 'x3', 'x4', 'y']
       df = df.iloc[50:150]
       df['y'] = df['y'].apply(lambda x: 0 if x == 'Iris-versicolor' else 1)
       df.tail()
```

```
[275]:
```

| | x1 | x2 | x3 | x4 | y |
|-----|-----|-----|-----|-----|---|
| 145 | 6.7 | 3.0 | 5.2 | 2.3 | 1 |
| 146 | 6.3 | 2.5 | 5.0 | 1.9 | 1 |
| 147 | 6.5 | 3.0 | 5.2 | 2.0 | 1 |
| 148 | 6.2 | 3.4 | 5.4 | 2.3 | 1 |

149 5.9 3.0 5.1 1.8 1

[200]: *# Assign features and target*

```
X = torch.tensor(df[['x2', 'x3']].values, dtype=torch.float)
y = torch.tensor(df['y'].values, dtype=torch.int)
```

Shuffling & train/test split

```
torch.manual_seed(123)
shuffle_idx = torch.randperm(y.size(0), dtype=torch.long)
```

```
X, y = X[shuffle_idx], y[shuffle_idx]
```

```
percent70 = int(shuffle_idx.size(0)*0.7)
```


```
X_train, X_test = X[shuffle_idx[:percent70]], X[shuffle_idx[percent70:]]
y_train, y_test = y[shuffle_idx[:percent70]], y[shuffle_idx[percent70:]]
```

Normalize (mean zero, unit variance)


```
mu, sigma = X_train.mean(dim=0), X_train.std(dim=0)
X_train = (X_train - mu) / sigma
X_test = (X_test - mu) / sigma
```

[202]:

```
plt.scatter(X_train[y_train == 0, 0], X_train[y_train == 0, 1], label='class 0')
plt.scatter(X_train[y_train == 1, 0], X_train[y_train == 1, 1], label='class 1')
plt.legend()
plt.show()
```



```
[203]: plt.scatter(X_test[y_test == 0, 0], X_test[y_test == 0, 1], label='class 0')
plt.scatter(X_test[y_test == 1, 0], X_test[y_test == 1, 1], label='class 1')
plt.legend()
plt.show()
```



```
[204]: # implement ADALINE
class Adaline1():
    def __init__(self, num_features):
        self.num_features = num_features
        self.weights = torch.zeros(num_features, 1, dtype=torch.float)
        self.bias = torch.zeros(1, dtype=torch.float)

    def forward(self, x):
        netinputs = torch.add(torch.mm(x, self.weights), self.bias)
        activations = netinputs
        return activations.view(-1)

    def backward(self, x, yhat, y):

        grad_loss_yhat = 2*(yhat - y)

        grad_yhat_weights = x
        grad_yhat_bias = 1.

        # Chain rule: inner times outer
        grad_loss_weights = torch.mm(grad_yhat_weights.t(), grad_loss_yhat.
↪view(-1, 1)) / y.size(0)

        grad_loss_bias = torch.sum(grad_yhat_bias*grad_loss_yhat) / y.size(0)

        # return negative gradient
        return (-1)*grad_loss_weights, (-1)*grad_loss_bias
```

```
[205]: def loss(yhat, y):
        return torch.mean((yhat - y)**2)

def train(model, x, y, num_epochs,
          learning_rate=0.01, seed=123, minibatch_size=10):
    cost = []

    torch.manual_seed(seed)
    for e in range(num_epochs):

        ##### Shuffle epoch
        shuffle_idx = torch.randperm(y.size(0), dtype=torch.long)
        minibatches = torch.split(shuffle_idx, minibatch_size)

        for minibatch_idx in minibatches:

            ##### Compute outputs #####
            yhat = model.forward(x[minibatch_idx])
```

```

##### Compute gradients #####
negative_grad_w, negative_grad_b = \
    model.backward(x[minibatch_idx], yhat, y[minibatch_idx])

##### Update weights #####
model.weights += learning_rate * negative_grad_w
model.bias += learning_rate * negative_grad_b

##### Logging #####
minibatch_loss = loss(yhat, y[minibatch_idx])
print('    Minibatch MSE: %.3f' % minibatch_loss)

##### Logging #####
yhat = model.forward(x)
curr_loss = loss(yhat, y)
print('Epoch: %03d' % (e+1), end="")
print(' | MSE: %.5f' % curr_loss)
cost.append(curr_loss)

return cost

```

```

[206]: model = Adaline1(num_features=X_train.size(1))
cost = train(model, X_train, y_train.float(), num_epochs=20, learning_rate=0.
↪1, seed=123, minibatch_size=10)

```

```

Minibatch MSE: 0.500
Minibatch MSE: 0.341
Minibatch MSE: 0.220
Minibatch MSE: 0.245
Minibatch MSE: 0.157
Minibatch MSE: 0.133
Minibatch MSE: 0.144
Epoch: 001 | MSE: 0.12142
Minibatch MSE: 0.107
Minibatch MSE: 0.147
Minibatch MSE: 0.064
Minibatch MSE: 0.079
Minibatch MSE: 0.185
Minibatch MSE: 0.063
Minibatch MSE: 0.135
Epoch: 002 | MSE: 0.09932
Minibatch MSE: 0.093
Minibatch MSE: 0.064
Minibatch MSE: 0.128
Minibatch MSE: 0.099
Minibatch MSE: 0.079

```


Minibatch MSE: 0.157
Minibatch MSE: 0.080
Epoch: 003 | MSE: 0.09693
Minibatch MSE: 0.131
Minibatch MSE: 0.146
Minibatch MSE: 0.050
Minibatch MSE: 0.095
Minibatch MSE: 0.106
Minibatch MSE: 0.072
Minibatch MSE: 0.102
Epoch: 004 | MSE: 0.09658
Minibatch MSE: 0.107
Minibatch MSE: 0.204
Minibatch MSE: 0.149
Minibatch MSE: 0.054
Minibatch MSE: 0.060
Minibatch MSE: 0.056
Minibatch MSE: 0.069
Epoch: 005 | MSE: 0.09657
Minibatch MSE: 0.068
Minibatch MSE: 0.111
Minibatch MSE: 0.092
Minibatch MSE: 0.115
Minibatch MSE: 0.157
Minibatch MSE: 0.074
Minibatch MSE: 0.087
Epoch: 006 | MSE: 0.09650
Minibatch MSE: 0.057
Minibatch MSE: 0.070
Minibatch MSE: 0.133
Minibatch MSE: 0.127
Minibatch MSE: 0.062
Minibatch MSE: 0.153
Minibatch MSE: 0.103
Epoch: 007 | MSE: 0.09683
Minibatch MSE: 0.102
Minibatch MSE: 0.110
Minibatch MSE: 0.101
Minibatch MSE: 0.065
Minibatch MSE: 0.126
Minibatch MSE: 0.124
Minibatch MSE: 0.076
Epoch: 008 | MSE: 0.09648
Minibatch MSE: 0.120
Minibatch MSE: 0.056
Minibatch MSE: 0.100
Minibatch MSE: 0.102
Minibatch MSE: 0.106

Minibatch MSE: 0.075
Minibatch MSE: 0.144
Epoch: 009 | MSE: 0.09740
Minibatch MSE: 0.073
Minibatch MSE: 0.071
Minibatch MSE: 0.084
Minibatch MSE: 0.152
Minibatch MSE: 0.099
Minibatch MSE: 0.108
Minibatch MSE: 0.118
Epoch: 010 | MSE: 0.09636
Minibatch MSE: 0.058
Minibatch MSE: 0.070
Minibatch MSE: 0.145
Minibatch MSE: 0.081
Minibatch MSE: 0.093
Minibatch MSE: 0.127
Minibatch MSE: 0.115
Epoch: 011 | MSE: 0.09638
Minibatch MSE: 0.123
Minibatch MSE: 0.091
Minibatch MSE: 0.085
Minibatch MSE: 0.093
Minibatch MSE: 0.091
Minibatch MSE: 0.143
Minibatch MSE: 0.081
Epoch: 012 | MSE: 0.09718
Minibatch MSE: 0.096
Minibatch MSE: 0.076
Minibatch MSE: 0.149
Minibatch MSE: 0.092
Minibatch MSE: 0.116
Minibatch MSE: 0.093
Minibatch MSE: 0.091
Epoch: 013 | MSE: 0.09638
Minibatch MSE: 0.095
Minibatch MSE: 0.104
Minibatch MSE: 0.107
Minibatch MSE: 0.120
Minibatch MSE: 0.102
Minibatch MSE: 0.045
Minibatch MSE: 0.124
Epoch: 014 | MSE: 0.09685
Minibatch MSE: 0.121
Minibatch MSE: 0.051
Minibatch MSE: 0.095
Minibatch MSE: 0.122
Minibatch MSE: 0.030


```
Minibatch MSE: 0.158
Minibatch MSE: 0.121
Epoch: 015 | MSE: 0.09745
Minibatch MSE: 0.080
Minibatch MSE: 0.119
Minibatch MSE: 0.091
Minibatch MSE: 0.095
Minibatch MSE: 0.044
Minibatch MSE: 0.092
Minibatch MSE: 0.180
Epoch: 016 | MSE: 0.09693
Minibatch MSE: 0.054
Minibatch MSE: 0.075
Minibatch MSE: 0.184
Minibatch MSE: 0.105
Minibatch MSE: 0.121
Minibatch MSE: 0.066
Minibatch MSE: 0.096
Epoch: 017 | MSE: 0.09699
Minibatch MSE: 0.076
Minibatch MSE: 0.050
Minibatch MSE: 0.198
Minibatch MSE: 0.105
Minibatch MSE: 0.054
Minibatch MSE: 0.136
Minibatch MSE: 0.099
Epoch: 018 | MSE: 0.09672
Minibatch MSE: 0.158
Minibatch MSE: 0.099
Minibatch MSE: 0.087
Minibatch MSE: 0.070
Minibatch MSE: 0.103
Minibatch MSE: 0.112
Minibatch MSE: 0.081
Epoch: 019 | MSE: 0.09638
Minibatch MSE: 0.095
Minibatch MSE: 0.093
Minibatch MSE: 0.111
Minibatch MSE: 0.147
Minibatch MSE: 0.083
Minibatch MSE: 0.102
Minibatch MSE: 0.065
Epoch: 020 | MSE: 0.09635
```

```
[208]: plt.plot(range(len(cost)), cost)
plt.ylabel('Mean Squared Error')
plt.xlabel('Epoch')
```

```
plt.show()
```



```
[209]: ones = torch.ones(y_train.size())
zeros = torch.zeros(y_train.size())
train_pred = model.forward(X_train)
train_acc = torch.mean(
    (torch.where(train_pred > 0.5,
                 ones,
                 zeros).int() == y_train).float())

ones = torch.ones(y_test.size())
zeros = torch.zeros(y_test.size())
test_pred = model.forward(X_test)
test_acc = torch.mean(
    (torch.where(test_pred > 0.5,
                 ones,
                 zeros).int() == y_test).float())

print('Training Accuracy: %.2f' % (train_acc*100))
print('Test Accuracy: %.2f' % (test_acc*100))
```

Training Accuracy: 90.00

Test Accuracy: 96.67

0.5 Automatic differentiation with PyTorch

```
[210]: import torch
from torch.autograd import grad
import torch.nn.functional as F
```

```
[212]: x = torch.tensor([3.])
w = torch.tensor([2.], requires_grad=True)
b = torch.tensor([1.], requires_grad=True)
a = F.relu(x*w + b)
```

```
[213]: a
```

```
[213]: tensor([7.], grad_fn=<ReluBackward0>)
```

```
[214]: grad(a, w, retain_graph=True)
```

```
[214]: (tensor([3.]),)
```

```
[215]: grad(a, b)
```

```
[215]: (tensor([1.]),)
```

```
[216]: grad(a, b)
```

```
-----
RuntimeError                                Traceback (most recent call last)
```

```
Input In [216], in <cell line: 1>()
```

```
----> 1 grad(a, b)
```

```
File /opt/anaconda3/lib/python3.9/site-packages/torch/autograd/__init__.py:300,
```

```
↳ in grad(outputs, inputs, grad_outputs, retain_graph, create_graph, ↵
```

```
↳ only_inputs, allow_unused, is_grads_batched)
```

```
    298     return vmap_internals._vmap(vjp, 0, 0, ↵
```

```
↳ allow_none_pass_through=True)(grad_outputs_)
```

```
    299 else:
```

```
--> 300     return ↵
```

```
↳ Variable._execution_engine.run_backward( # Calls into the C++ engine to run the backward ↵
```

```
    301         t_outputs, grad_outputs_, retain_graph, create_graph, t_inputs, ↵
```

```
    302         allow_unused, accumulate_grad=False)
```

```
RuntimeError: Trying to backward through the graph a second time (or directly ↵
```

```
↳ access saved tensors after they have already been freed). Saved intermediate ↵
```

```
↳ values of the graph are freed when you call .backward() or autograd.grad(). ↵
```

```
↳ Specify retain_graph=True if you need to backward through the graph a second ↵
```

```
↳ time or if you need to access saved tensors after calling backward.
```

```
[217]: x = torch.tensor([3.])
w = torch.tensor([2.], requires_grad=True)
b = torch.tensor([1.], requires_grad=True)

def my_relu(z):
    if z > 0.:
        return z
    else:
        z[:] = 0.
        return z

a = my_relu(x*w + b)
grad(a, w)
```

```
[217]: (tensor([3.]),)
```

0.5.1 What happens to ReLU function at 0 ?

```
[218]: x = torch.tensor([-1.])
w = torch.tensor([1.], requires_grad=True)
b = torch.tensor([1.], requires_grad=True)

def my_relu(z):
    if z > 0.:
        return z
    else:
        z[:] = 0.
        return z

a = F.relu(x*w + b)
grad(a, w, retain_graph=False)
```

```
[218]: (tensor([-0.]),)
```

0.6 Adaline With Pytorch

```
[247]: df = pd.read_csv('data/iris.data', index_col=None, header=None)
df.columns = ['x1', 'x2', 'x3', 'x4', 'y']
df = df.iloc[50:150]
df['y'] = df['y'].apply(lambda x: 0 if x == 'Iris-versicolor' else 1)

# Assign features and target

X = torch.tensor(df[['x2', 'x3']].values, dtype=torch.float)
```

```

y = torch.tensor(df['y'].values, dtype=torch.int)

# Shuffling & train/test split

torch.manual_seed(123)
shuffle_idx = torch.randperm(y.size(0), dtype=torch.long)

X, y = X[shuffle_idx], y[shuffle_idx]

percent70 = int(shuffle_idx.size(0)*0.7)

X_train, X_test = X[shuffle_idx[:percent70]], X[shuffle_idx[percent70:]]
y_train, y_test = y[shuffle_idx[:percent70]], y[shuffle_idx[percent70:]]

# Normalize (mean zero, unit variance)

mu, sigma = X_train.mean(dim=0), X_train.std(dim=0)
X_train = (X_train - mu) / sigma
X_test = (X_test - mu) / sigma

```

```

[221]: class Adaline2():
    def __init__(self, num_features):
        self.num_features = num_features
        self.weight = torch.zeros(num_features, 1, dtype=torch.
↪float, requires_grad=True)
        self.bias = torch.zeros(1, dtype=torch.float, requires_grad=True)

    def forward(self, x):
        netinputs = torch.add(torch.mm(x, self.weight), self.bias)
        activations = netinputs
        return activations.view(-1)

def loss_func(yhat, y):
    return torch.mean((yhat - y)**2)

def train(model, x, y, num_epochs, learning_rate=0.01, seed=123, ↪
↪minibatch_size=10):
    cost = []

    torch.manual_seed(seed)
    for e in range(num_epochs):

        #### Shuffle epoch
        shuffle_idx = torch.randperm(y.size(0), dtype=torch.long)
        minibatches = torch.split(shuffle_idx, minibatch_size)

```

```

for minibatch_idx in minibatches:

    ##### Compute outputs #####
    yhat = model.forward(x[minibatch_idx])
    loss = loss_func(yhat, y[minibatch_idx])

    ##### Compute gradients #####
    negative_grad_w = grad(loss, model.weight, retain_graph=True)[0] * (-1)

    negative_grad_b = grad(loss, model.bias)[0] * (-1)

    ##### Update weights #####
    model.weight = model.weight + learning_rate * negative_grad_w
    model.bias = model.bias + learning_rate * negative_grad_b

    ##### Logging #####
    with torch.no_grad():
        # context manager to
        # avoid building graph during "inference"
        # to save memory
        yhat = model.forward(x)
        curr_loss = loss_func(yhat, y)
        print('Epoch: %03d' % (e+1), end="")
        print(' | MSE: %.5f' % curr_loss)
        cost.append(curr_loss)

return cost

```

```

[234]: model = Adaline2(num_features=X_train.size(1))
cost = train(model,
              X_train, y_train.float(),
              num_epochs=20,
              learning_rate=0.01,
              seed=123,
              minibatch_size=10)

```


```

Epoch: 001 | MSE: 0.38849
Epoch: 002 | MSE: 0.31679
Epoch: 003 | MSE: 0.26379
Epoch: 004 | MSE: 0.22463
Epoch: 005 | MSE: 0.19527
Epoch: 006 | MSE: 0.17307
Epoch: 007 | MSE: 0.15629
Epoch: 008 | MSE: 0.14352
Epoch: 009 | MSE: 0.13360

```

```
Epoch: 010 | MSE: 0.12600
Epoch: 011 | MSE: 0.12007
Epoch: 012 | MSE: 0.11547
Epoch: 013 | MSE: 0.11178
Epoch: 014 | MSE: 0.10884
Epoch: 015 | MSE: 0.10656
Epoch: 016 | MSE: 0.10470
Epoch: 017 | MSE: 0.10320
Epoch: 018 | MSE: 0.10200
Epoch: 019 | MSE: 0.10105
Epoch: 020 | MSE: 0.10025
```

```
[235]: plt.plot(range(len(cost)), cost)
plt.ylabel('Mean Squared Error')
plt.xlabel('Epoch')
plt.show()
```



```
[236]: ones = torch.ones(y_train.size())
zeros = torch.zeros(y_train.size())
train_pred = model.forward(X_train)
train_acc = torch.mean(
    (torch.where(train_pred > 0.5,
                 ones,
                 zeros).int() == y_train).float())
```

```

ones = torch.ones(y_test.size())
zeros = torch.zeros(y_test.size())
test_pred = model.forward(X_test)
test_acc = torch.mean(
    (torch.where(test_pred > 0.5,
                 ones,
                 zeros).int() == y_test).float())

print('Training Accuracy: %.2f' % (train_acc*100))
print('Test Accuracy: %.2f' % (test_acc*100))

```

Training Accuracy: 92.86
Test Accuracy: 93.33

0.7 Adaline With Pytorch With Automatic Diff

```

[260]: class Adaline3(torch.nn.Module):
    def __init__(self, num_features):
        super(Adaline3, self).__init__()
        self.linear = torch.nn.Linear(num_features, 1)

        self.linear.weight.detach().zero_()
        self.linear.bias.detach().zero_()

    def forward(self, x):
        netinputs = self.linear(x)
        activations = netinputs
        return activations.view(-1)

def train(model, x, y, num_epochs, learning_rate = 0.01, seed = 123,
    minibatch_size=10):
    cost = []
    torch.manual_seed(seed)
    optimiser = torch.optim.SGD(model.parameters(), lr=learning_rate)

    for e in range(num_epochs):
        shuffle_idx = torch.randperm(y.size(0), dtype=torch.long)
        minibatches = torch.split(shuffle_idx, minibatch_size)

        for minibatch_idx in minibatches:

            yhat = model.forward(x[minibatch_idx])
            loss = F.mse_loss(yhat, y[minibatch_idx])

            optimiser.zero_grad()

```



```

        loss.backward()

        optimiser.step()

    with torch.no_grad():
        yhat = model.forward(x)
        curr_loss = F.mse_loss(yhat, y)

        print('Epoch: %03d' % (e+1), end="")
        print(' | MSE: %.5f' % curr_loss)
        cost.append(curr_loss)

    return cost;

```

```

[261]: model = Adaline3(num_features=X_train.size(1))
cost = train(model, X_train, y_train.float(), num_epochs=20, learning_rate=0.
↪01, seed=123, minibatch_size=10)

```

```


Epoch: 001 | MSE: 0.38849
Epoch: 002 | MSE: 0.31679
Epoch: 003 | MSE: 0.26379
Epoch: 004 | MSE: 0.22463
Epoch: 005 | MSE: 0.19527
Epoch: 006 | MSE: 0.17307
Epoch: 007 | MSE: 0.15629
Epoch: 008 | MSE: 0.14352
Epoch: 009 | MSE: 0.13360
Epoch: 010 | MSE: 0.12600
Epoch: 011 | MSE: 0.12007
Epoch: 012 | MSE: 0.11547
Epoch: 013 | MSE: 0.11178
Epoch: 014 | MSE: 0.10884
Epoch: 015 | MSE: 0.10656
Epoch: 016 | MSE: 0.10470
Epoch: 017 | MSE: 0.10320
Epoch: 018 | MSE: 0.10200
Epoch: 019 | MSE: 0.10105
Epoch: 020 | MSE: 0.10025

```

```

[262]: plt.plot(range(len(cost)), cost)
plt.ylabel('Mean Squared Error')
plt.xlabel('Epoch')
plt.show()

```



[]: