

## lab-6-softmax-regression

March 4, 2023

```
[1]: %matplotlib inline
import matplotlib.pyplot as plt
import pandas as pd
import torch
import torch.nn.functional as F
```

```
[14]: df = pd.read_csv('./data/iris.data', index_col=None, header=None)
df.columns = ['x1', 'x2', 'x3', 'x4', 'y']

d = {'Iris-versicolor': 1,
      'Iris-virginica': 2,
      'Iris-setosa': 0,
    }

df['y'] = df['y'].map(d)

# Assign features and target

X = torch.tensor(df[['x2', 'x4']].values, dtype=torch.float)
y = torch.tensor(df['y'].values, dtype=torch.int)

# Shuffling & train/test split

torch.manual_seed(123)
shuffle_idx = torch.randperm(y.size(0), dtype=torch.long)

X, y = X[shuffle_idx], y[shuffle_idx]

percent80 = int(shuffle_idx.size(0)*0.8)


X_train, X_test = X[shuffle_idx[:percent80]], X[shuffle_idx[percent80:]]
y_train, y_test = y[shuffle_idx[:percent80]], y[shuffle_idx[percent80:]]

# Normalize (mean zero, unit variance)

mu, sigma = X_train.mean(dim=0), X_train.std(dim=0)
X_train = (X_train - mu) / sigma
```

```
X_test = (X_test - mu) / sigma
```

```
fig, ax = plt.subplots(1, 2, figsize=(7, 2.5))
ax[0].scatter(X_train[y_train == 2, 0], X_train[y_train == 2, 1])
ax[0].scatter(X_train[y_train == 1, 0], X_train[y_train == 1, 1], marker='v')
ax[0].scatter(X_train[y_train == 0, 0], X_train[y_train == 0, 1], marker='s')
ax[1].scatter(X_test[y_test == 2, 0], X_test[y_test == 2, 1])
ax[1].scatter(X_test[y_test == 1, 0], X_test[y_test == 1, 1], marker='v')
ax[1].scatter(X_test[y_test == 0, 0], X_test[y_test == 0, 1], marker='s')
plt.show()
```



## High Level Implementation

```
[6]: DEVICE = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
```

```
[23]: class SoftmaxRegression(torch.nn.Module):
    def __init__(self, num_features, num_class):
        super(SoftmaxRegression, self).__init__()
        self.linear = torch.nn.Linear(num_features, num_class, dtype=torch.
↪float32, device=DEVICE)

        self.linear.weight.detach().zero_()
        self.linear.bias.detach().zero_()
    def forward(self, x):
        logits = self.linear(x);
        probas = F.softmax(logits, dim=1)
        return logits, probas

    def predict_labels(self, x):
        logits, probas = self.forward(x)
        labels = torch.argmax(probas, dim=1)
```

```

        return labels

    def evaluate(self, x, y):
        labels = self.predict_labels(x).float()
        accuracy = torch.sum(labels.view(-1) == y.float()).item() / y.size(0)
        return accuracy

model = SoftmaxRegression(num_features=2, num_class=3).to(DEVICE)
optimizer = torch.optim.SGD(model.parameters(), lr=0.1)

```

```

[24]: def comp_accuracy(true_labels, pred_labels):
        accuracy = torch.sum(true_labels.view(-1).float() ==
                               pred_labels.float()).item() / true_labels.size(0)
        return accuracy

X_train = X_train.to(DEVICE)
y_train = y_train.to(DEVICE)
X_test = X_test.to(DEVICE)
y_test = y_test.to(DEVICE)

num_epochs = 50
for epoch in range(num_epochs):

    logits, probas = model(X_train)

    # Compute gradients
    cost = F.cross_entropy(probas, y_train.long())
    optimser.zero_grad()
    cost.backward();

    # Update weights
    optimizer.step()

    logits, probas = model(X_train)
    acc = comp_accuracy(y_train, torch.argmax(probas, dim=1))
    print('Epoch: %03d' % (epoch + 1), end="")
    print(' | Train ACC: %.3f' % acc, end="")
    print(' | Cost: %.3f' % F.cross_entropy(logits, y_train.long()))

print('\nModel parameters:')
print('  Weights: %s' % model.linear.weight)
print('  Bias: %s' % model.linear.bias)

```

Epoch: 001 | Train ACC: 0.342 | Cost: 1.199

Epoch: 002		Train ACC: 0.342		Cost: 1.167
Epoch: 003		Train ACC: 0.342		Cost: 1.122
Epoch: 004		Train ACC: 0.342		Cost: 1.065
Epoch: 005		Train ACC: 0.342		Cost: 1.000
Epoch: 006		Train ACC: 0.358		Cost: 0.929
Epoch: 007		Train ACC: 0.542		Cost: 0.856
Epoch: 008		Train ACC: 0.658		Cost: 0.784
Epoch: 009		Train ACC: 0.708		Cost: 0.716
Epoch: 010		Train ACC: 0.767		Cost: 0.653
Epoch: 011		Train ACC: 0.767		Cost: 0.597
Epoch: 012		Train ACC: 0.783		Cost: 0.549
Epoch: 013		Train ACC: 0.783		Cost: 0.507
Epoch: 014		Train ACC: 0.792		Cost: 0.473
Epoch: 015		Train ACC: 0.783		Cost: 0.444
Epoch: 016		Train ACC: 0.783		Cost: 0.421
Epoch: 017		Train ACC: 0.783		Cost: 0.403
Epoch: 018		Train ACC: 0.792		Cost: 0.388
Epoch: 019		Train ACC: 0.800		Cost: 0.376
Epoch: 020		Train ACC: 0.792		Cost: 0.367
Epoch: 021		Train ACC: 0.800		Cost: 0.360
Epoch: 022		Train ACC: 0.800		Cost: 0.355
Epoch: 023		Train ACC: 0.800		Cost: 0.351
Epoch: 024		Train ACC: 0.825		Cost: 0.348
Epoch: 025		Train ACC: 0.833		Cost: 0.346
Epoch: 026		Train ACC: 0.833		Cost: 0.344
Epoch: 027		Train ACC: 0.833		Cost: 0.344
Epoch: 028		Train ACC: 0.833		Cost: 0.343
Epoch: 029		Train ACC: 0.842		Cost: 0.344
Epoch: 030		Train ACC: 0.867		Cost: 0.345
Epoch: 031		Train ACC: 0.858		Cost: 0.346
Epoch: 032		Train ACC: 0.858		Cost: 0.347
Epoch: 033		Train ACC: 0.858		Cost: 0.349
Epoch: 034		Train ACC: 0.858		Cost: 0.352
Epoch: 035		Train ACC: 0.858		Cost: 0.355
Epoch: 036		Train ACC: 0.858		Cost: 0.358
Epoch: 037		Train ACC: 0.883		Cost: 0.361
Epoch: 038		Train ACC: 0.883		Cost: 0.365
Epoch: 039		Train ACC: 0.883		Cost: 0.369
Epoch: 040		Train ACC: 0.883		Cost: 0.373
Epoch: 041		Train ACC: 0.883		Cost: 0.378
Epoch: 042		Train ACC: 0.883		Cost: 0.383
Epoch: 043		Train ACC: 0.883		Cost: 0.388
Epoch: 044		Train ACC: 0.883		Cost: 0.394
Epoch: 045		Train ACC: 0.875		Cost: 0.400
Epoch: 046		Train ACC: 0.883		Cost: 0.406
Epoch: 047		Train ACC: 0.883		Cost: 0.412
Epoch: 048		Train ACC: 0.883		Cost: 0.419
Epoch: 049		Train ACC: 0.883		Cost: 0.425

Epoch: 050 | Train ACC: 0.883 | Cost: 0.432

Model parameters:

Weights: Parameter containing:  
tensor([[ 4.0129, -7.3231],  
 [-4.1433, 0.0976],  
 [ 0.1304, 7.2254]], requires\_grad=True)  
Bias: Parameter containing:  
tensor([-2.0453, 1.4199, -0.1037], requires\_grad=True)

```
[25]: X_test = X_test.to(DEVICE)
      y_test = y_test.to(DEVICE)

      test_acc = model.evaluate(X_test, y_test)
      print('Test set accuracy: %.2f%%' % (test_acc*100))
```

Test set accuracy: 83.33%

## 0.1 Softmax Regression on MNIST

```
[61]: from torchvision import datasets
      from torchvision import transforms
      from torch.utils.data import DataLoader
      import time
```

```
[62]: # Device
      device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")

      # Hyperparameters
      random_seed = 123
      learning_rate = 0.1
      num_epochs = 25
      batch_size = 256

      # Architecture
      num_features = 784
      num_classes = 10

      train_dataset = datasets.MNIST(root='data',
                                     train=True,
                                     transform=transforms.ToTensor(),
                                     download=True)

      test_dataset = datasets.MNIST(root='data',
                                    train=False,
                                    transform=transforms.ToTensor())
```

```

train_loader = DataLoader(dataset=train_dataset,
                           batch_size=batch_size,
                           shuffle=True)

test_loader = DataLoader(dataset=test_dataset,
                          batch_size=batch_size,
                          shuffle=False)

# Checking the dataset
for images, labels in train_loader:
    print('Image batch dimensions:', images.shape) #NCHW
    print('Image label dimensions:', labels.shape)
    break

```

Image batch dimensions: torch.Size([256, 1, 28, 28])  
Image label dimensions: torch.Size([256])

```
[63]: labels[:10]
```

```
[63]: tensor([0, 8, 8, 4, 3, 7, 6, 3, 2, 5])
```

```

[64]: class SoftmaxRegression(torch.nn.Module):

    def __init__(self, num_features, num_classes):
        super(SoftmaxRegression, self).__init__()
        self.linear = torch.nn.Linear(num_features, num_classes)

        self.linear.weight.detach().zero_()
        self.linear.bias.detach().zero_()

    def forward(self, x):
        logits = self.linear(x)
        probas = F.softmax(logits, dim=1)
        return logits, probas

model = SoftmaxRegression(num_features=num_features,
                           num_classes=num_classes)

model.to(device)

```

```

[64]: SoftmaxRegression(
  (linear): Linear(in_features=784, out_features=10, bias=True)
)

```

```
[65]: optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)
```

```
[66]: # Manual seed for deterministic data loader
torch.manual_seed(random_seed)

def compute_accuracy(model, data_loader):
    correct_pred, num_examples = 0, 0

    for features, targets in data_loader:
        features = features.view(-1, 28*28).to(device)
        targets = targets.to(device)
        logits, probas = model(features)
        _, predicted_labels = torch.max(probas, 1)
        num_examples += targets.size(0)
        correct_pred += (predicted_labels == targets).sum()

    return correct_pred.float() / num_examples * 100
```

```
[67]: for f in train_loader:
        print(f[0].shape)
        print(f[1].shape)
        break;
```

```
torch.Size([256, 1, 28, 28])
torch.Size([256])
```

```
[68]: start_time = time.time()
epoch_costs = []
for epoch in range(num_epochs):
    avg_cost = 0.
    for batch_idx, (features, targets) in enumerate(train_loader):
        features = features.view(-1, 28*28).to(device)
        targets = targets.to(device)

        # forward propagation
        logits, probas = model(features)

        # calculate loss
        cost = F.cross_entropy(logits, targets)
        # calculate gradient
        optimizer.zero_grad()
        cost.backward()
        avg_cost += cost

        # update model params
        optimizer.step()
```

```

# log
if not batch_idx % 50:
    print ('Epoch: %03d/%03d | Batch %03d/%03d | Cost: %.4f'
           %(epoch+1, num_epochs, batch_idx,
             len(train_dataset)//batch_size, cost))

with torch.set_grad_enabled(False):
    avg_cost = avg_cost/len(train_dataset)
    epoch_costs.append(avg_cost)
    print('Epoch: %03d/%03d training accuracy: %.2f%%' % (epoch+1,
    num_epochs, compute_accuracy(model, train_loader)))
    print('Time elapsed: %.2f min' % ((time.time() - start_time)/60))

```

```

Epoch: 001/025 | Batch 000/234 | Cost: 2.3026
Epoch: 001/025 | Batch 050/234 | Cost: 0.7769
Epoch: 001/025 | Batch 100/234 | Cost: 0.6825
Epoch: 001/025 | Batch 150/234 | Cost: 0.5251
Epoch: 001/025 | Batch 200/234 | Cost: 0.5382
Epoch: 001/025 training accuracy: 87.92%
Time elapsed: 0.12 min
Epoch: 002/025 | Batch 000/234 | Cost: 0.5015
Epoch: 002/025 | Batch 050/234 | Cost: 0.4031
Epoch: 002/025 | Batch 100/234 | Cost: 0.4447
Epoch: 002/025 | Batch 150/234 | Cost: 0.4757
Epoch: 002/025 | Batch 200/234 | Cost: 0.4474
Epoch: 002/025 training accuracy: 89.28%
Time elapsed: 0.23 min
Epoch: 003/025 | Batch 000/234 | Cost: 0.4091
Epoch: 003/025 | Batch 050/234 | Cost: 0.4192
Epoch: 003/025 | Batch 100/234 | Cost: 0.2462
Epoch: 003/025 | Batch 150/234 | Cost: 0.4043
Epoch: 003/025 | Batch 200/234 | Cost: 0.3774
Epoch: 003/025 training accuracy: 89.82%
Time elapsed: 0.35 min
Epoch: 004/025 | Batch 000/234 | Cost: 0.3269
Epoch: 004/025 | Batch 050/234 | Cost: 0.3283
Epoch: 004/025 | Batch 100/234 | Cost: 0.3350
Epoch: 004/025 | Batch 150/234 | Cost: 0.3011
Epoch: 004/025 | Batch 200/234 | Cost: 0.4058
Epoch: 004/025 training accuracy: 90.35%
Time elapsed: 0.46 min
Epoch: 005/025 | Batch 000/234 | Cost: 0.3713
Epoch: 005/025 | Batch 050/234 | Cost: 0.3954
Epoch: 005/025 | Batch 100/234 | Cost: 0.3983

```




Epoch: 005/025 | Batch 150/234 | Cost: 0.2824  
Epoch: 005/025 | Batch 200/234 | Cost: 0.3520  
Epoch: 005/025 training accuracy: 90.62%  
Time elapsed: 0.59 min  
Epoch: 006/025 | Batch 000/234 | Cost: 0.2830  
Epoch: 006/025 | Batch 050/234 | Cost: 0.3104  
Epoch: 006/025 | Batch 100/234 | Cost: 0.3334  
Epoch: 006/025 | Batch 150/234 | Cost: 0.3317  
Epoch: 006/025 | Batch 200/234 | Cost: 0.3608  
Epoch: 006/025 training accuracy: 90.82%  
Time elapsed: 0.71 min  
Epoch: 007/025 | Batch 000/234 | Cost: 0.3388  
Epoch: 007/025 | Batch 050/234 | Cost: 0.3077  
Epoch: 007/025 | Batch 100/234 | Cost: 0.3147  
Epoch: 007/025 | Batch 150/234 | Cost: 0.2924  
Epoch: 007/025 | Batch 200/234 | Cost: 0.3148  
Epoch: 007/025 training accuracy: 91.03%  
Time elapsed: 0.84 min  
Epoch: 008/025 | Batch 000/234 | Cost: 0.3366  
Epoch: 008/025 | Batch 050/234 | Cost: 0.3519  
Epoch: 008/025 | Batch 100/234 | Cost: 0.3838  
Epoch: 008/025 | Batch 150/234 | Cost: 0.2811  
Epoch: 008/025 | Batch 200/234 | Cost: 0.3233  
Epoch: 008/025 training accuracy: 91.17%  
Time elapsed: 0.97 min  
Epoch: 009/025 | Batch 000/234 | Cost: 0.3393  
Epoch: 009/025 | Batch 050/234 | Cost: 0.3493  
Epoch: 009/025 | Batch 100/234 | Cost: 0.3406  
Epoch: 009/025 | Batch 150/234 | Cost: 0.3271  
Epoch: 009/025 | Batch 200/234 | Cost: 0.2886  
Epoch: 009/025 training accuracy: 91.29%  
Time elapsed: 1.10 min  
Epoch: 010/025 | Batch 000/234 | Cost: 0.2683  
Epoch: 010/025 | Batch 050/234 | Cost: 0.2920  
Epoch: 010/025 | Batch 100/234 | Cost: 0.2856  
Epoch: 010/025 | Batch 150/234 | Cost: 0.3030  
Epoch: 010/025 | Batch 200/234 | Cost: 0.3434  
Epoch: 010/025 training accuracy: 91.42%  
Time elapsed: 1.23 min  
Epoch: 011/025 | Batch 000/234 | Cost: 0.3203  
Epoch: 011/025 | Batch 050/234 | Cost: 0.2722  
Epoch: 011/025 | Batch 100/234 | Cost: 0.3426  
Epoch: 011/025 | Batch 150/234 | Cost: 0.3474  
Epoch: 011/025 | Batch 200/234 | Cost: 0.2716  
Epoch: 011/025 training accuracy: 91.56%  
Time elapsed: 1.37 min  
Epoch: 012/025 | Batch 000/234 | Cost: 0.3234  
Epoch: 012/025 | Batch 050/234 | Cost: 0.3058

Epoch: 012/025 | Batch 100/234 | Cost: 0.3143  
Epoch: 012/025 | Batch 150/234 | Cost: 0.2685  
Epoch: 012/025 | Batch 200/234 | Cost: 0.2487  
Epoch: 012/025 training accuracy: 91.64%  
Time elapsed: 1.50 min  
Epoch: 013/025 | Batch 000/234 | Cost: 0.3649  
Epoch: 013/025 | Batch 050/234 | Cost: 0.2811  
Epoch: 013/025 | Batch 100/234 | Cost: 0.2619  
Epoch: 013/025 | Batch 150/234 | Cost: 0.3158  
Epoch: 013/025 | Batch 200/234 | Cost: 0.2522  
Epoch: 013/025 training accuracy: 91.68%  
Time elapsed: 1.64 min  
Epoch: 014/025 | Batch 000/234 | Cost: 0.3492  
Epoch: 014/025 | Batch 050/234 | Cost: 0.2866  
Epoch: 014/025 | Batch 100/234 | Cost: 0.2746  
Epoch: 014/025 | Batch 150/234 | Cost: 0.3830  
Epoch: 014/025 | Batch 200/234 | Cost: 0.3988  
Epoch: 014/025 training accuracy: 91.73%  
Time elapsed: 1.77 min  
Epoch: 015/025 | Batch 000/234 | Cost: 0.1983  
Epoch: 015/025 | Batch 050/234 | Cost: 0.3476  
Epoch: 015/025 | Batch 100/234 | Cost: 0.2963  
Epoch: 015/025 | Batch 150/234 | Cost: 0.2509  
Epoch: 015/025 | Batch 200/234 | Cost: 0.3132  
Epoch: 015/025 training accuracy: 91.83%  
Time elapsed: 1.91 min  
Epoch: 016/025 | Batch 000/234 | Cost: 0.2453  
Epoch: 016/025 | Batch 050/234 | Cost: 0.3091  
Epoch: 016/025 | Batch 100/234 | Cost: 0.2818  
Epoch: 016/025 | Batch 150/234 | Cost: 0.2680  
Epoch: 016/025 | Batch 200/234 | Cost: 0.2571  
Epoch: 016/025 training accuracy: 91.88%  
Time elapsed: 10.15 min  
Epoch: 017/025 | Batch 000/234 | Cost: 0.2611  
Epoch: 017/025 | Batch 050/234 | Cost: 0.3563  
Epoch: 017/025 | Batch 100/234 | Cost: 0.2167  
Epoch: 017/025 | Batch 150/234 | Cost: 0.3099  
Epoch: 017/025 | Batch 200/234 | Cost: 0.3305  
Epoch: 017/025 training accuracy: 91.86%  
Time elapsed: 10.27 min  
Epoch: 018/025 | Batch 000/234 | Cost: 0.2718  
Epoch: 018/025 | Batch 050/234 | Cost: 0.2744  
Epoch: 018/025 | Batch 100/234 | Cost: 0.3043  
Epoch: 018/025 | Batch 150/234 | Cost: 0.2605  
Epoch: 018/025 | Batch 200/234 | Cost: 0.2321  
Epoch: 018/025 training accuracy: 91.97%  
Time elapsed: 10.39 min  
Epoch: 019/025 | Batch 000/234 | Cost: 0.3145

Epoch: 019/025 | Batch 050/234 | Cost: 0.3303  
 Epoch: 019/025 | Batch 100/234 | Cost: 0.2645  
 Epoch: 019/025 | Batch 150/234 | Cost: 0.3165  
 Epoch: 019/025 | Batch 200/234 | Cost: 0.2818  
 Epoch: 019/025 training accuracy: 91.98%  
 Time elapsed: 10.50 min  
 Epoch: 020/025 | Batch 000/234 | Cost: 0.3309  
 Epoch: 020/025 | Batch 050/234 | Cost: 0.2844  
 Epoch: 020/025 | Batch 100/234 | Cost: 0.1714  
 Epoch: 020/025 | Batch 150/234 | Cost: 0.3025  
 Epoch: 020/025 | Batch 200/234 | Cost: 0.2996  
 Epoch: 020/025 training accuracy: 91.97%  
 Time elapsed: 10.62 min  
 Epoch: 021/025 | Batch 000/234 | Cost: 0.3753  
 Epoch: 021/025 | Batch 050/234 | Cost: 0.2494  
 Epoch: 021/025 | Batch 100/234 | Cost: 0.2973  
 Epoch: 021/025 | Batch 150/234 | Cost: 0.2705  
 Epoch: 021/025 | Batch 200/234 | Cost: 0.2683  
 Epoch: 021/025 training accuracy: 92.03%  
 Time elapsed: 10.74 min  
 Epoch: 022/025 | Batch 000/234 | Cost: 0.2227  
 Epoch: 022/025 | Batch 050/234 | Cost: 0.2263  
 Epoch: 022/025 | Batch 100/234 | Cost: 0.2010  
 Epoch: 022/025 | Batch 150/234 | Cost: 0.3033  
 Epoch: 022/025 | Batch 200/234 | Cost: 0.3440  
 Epoch: 022/025 training accuracy: 92.13%  
 Time elapsed: 10.85 min  
 Epoch: 023/025 | Batch 000/234 | Cost: 0.2959  
 Epoch: 023/025 | Batch 050/234 | Cost: 0.2255  
 Epoch: 023/025 | Batch 100/234 | Cost: 0.3101  
 Epoch: 023/025 | Batch 150/234 | Cost: 0.2969  
 Epoch: 023/025 | Batch 200/234 | Cost: 0.3026  
 Epoch: 023/025 training accuracy: 92.17%  
 Time elapsed: 10.97 min  
 Epoch: 024/025 | Batch 000/234 | Cost: 0.3101  
 Epoch: 024/025 | Batch 050/234 | Cost: 0.2736  
 Epoch: 024/025 | Batch 100/234 | Cost: 0.2523  
 Epoch: 024/025 | Batch 150/234 | Cost: 0.3125  
 Epoch: 024/025 | Batch 200/234 | Cost: 0.2340  
 Epoch: 024/025 training accuracy: 92.12%  
 Time elapsed: 11.08 min  
 Epoch: 025/025 | Batch 000/234 | Cost: 0.3074  
 Epoch: 025/025 | Batch 050/234 | Cost: 0.2794  
 Epoch: 025/025 | Batch 100/234 | Cost: 0.2232  
 Epoch: 025/025 | Batch 150/234 | Cost: 0.3031  
 Epoch: 025/025 | Batch 200/234 | Cost: 0.2244  
 Epoch: 025/025 training accuracy: 92.26%  
 Time elapsed: 11.20 min

```
[69]: %matplotlib inline
import matplotlib
import matplotlib.pyplot as plt

plt.plot(epoch_costs)
plt.ylabel('Avg Cross Entropy Loss\n(approximated by averaging over_\n↳minibatches)')
plt.xlabel('Epoch')
plt.show()
```



```
[70]: print('Test accuracy: %.2f%%' % (compute_accuracy(model, test_loader)))
```

Test accuracy: 92.23%

```
[ ]:
```